
EU AI Act Conformity Self-Assessment

TheHiveryIQ / Hive Civilization

Steve Rotzin, Founder

Version 1.0 — April 2026

Prepared under Regulation (EU) 2024/1689 — EU Artificial Intelligence Act

This document is a voluntary conformity self-assessment. It is not legal advice and does not constitute certification by a notified body. Enterprises should engage qualified EU AI Act counsel for deployment-specific compliance review.

Table of Contents

1. Executive Summary
2. System Classification
3. Article 12 — Record-Keeping and Logging
4. Article 13 — Transparency Obligations
5. Article 14 — Human Oversight
6. Article 22 — Automated Decision-Making (GDPR Art. 22 Overlap)
7. Technical Robustness — HiveTrust Risk Verification
8. Data Governance
9. Self-Assessment Conclusion
10. Live Endpoints for Auditors

1. Executive Summary

Hive Civilization operates an autonomous agent network — branded as HiveBank (financial settlement and transaction services) and HiveGate (agent onboarding, identity issuance, and access control) — in which software agents hold W3C Decentralized Identifiers (DIDs), execute machine-signed employment and liability contracts (HAHS 1.0.0), log every decision to a cryptographic Agent Transaction Graph, and settle disputes in USDC on the Base L2 blockchain.

This document constitutes a voluntary conformity self-assessment under the EU Artificial Intelligence Act (Regulation (EU) 2024/1689, in force from August 2024, with obligations phasing in through August 2026). It is prepared for the purpose of enterprise procurement due diligence, regulatory readiness review, and public trust signaling.¹

Classification summary: HiveBank / HiveGate services are assessed as Limited Risk AI systems under the EU AI Act, with specific obligations around transparency and logging that Hive meets through existing technical controls. Sections 3–8 document that conformance in detail.

This self-assessment does not constitute a formal notified-body certification. Hive Civilization commits to upgrading to third-party conformity assessment procedures if and when regulatory guidance or enterprise requirements mandate it.

1. EU AI Act full text: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj>

2. System Classification

2.1 Applicable Risk Category

Under the EU AI Act, AI systems are classified across four tiers: Unacceptable Risk (prohibited), High Risk (Annex III), Limited Risk, and Minimal Risk.

Classification Criterion	Assessment	Rationale
Prohibited practices (Art. 5)	Not applicable	Hive does not perform subliminal manipulation, social scoring, or real-time biometric surveillance
Annex III High-Risk categories	Not applicable at this time	HiveBank operates agent-to-agent financial settlement, not credit scoring of natural persons; HiveGate issues DIDs to agents, not to humans for access to essential services
General-purpose AI with systemic risk	Under review	If Hive's SDK is integrated into third-party deployments that touch Annex III use cases, downstream deployers bear classification responsibility
Limited Risk — AI systems that interact with natural persons	Applicable	HiveGate surfaces risk scores and identity confirmations that human compliance officers may act upon

Primary classification: Limited Risk AI System.

Hive acknowledges that if HiveBank's automated transaction logic is integrated into a credit, insurance, or employment workflow by a downstream deployer, that deployer's use case may independently qualify as High Risk under Annex III. Hive's API design supports that deployer's own conformity obligations through the transparency, logging, and explainability controls described below.

2.2 Scope Boundaries

Service	In Scope	Out of Scope
HiveBank — USDC settlement	Yes	Fiat on/off ramp (not operated by Hive)
HiveGate — DID issuance	Yes	Downstream use of DIDs by third-party agents
HiveTrust — risk verification API	Yes	Third-party integrations consuming the API
HAHS contract execution engine	Yes	Legal enforceability in specific jurisdictions (legal counsel review required)

3. Article 12 — Record-Keeping and Logging

EU AI Act Article 12 requires high-risk AI systems to automatically log events to the extent necessary to ensure post-hoc auditability.

Although Hive is classified as Limited Risk (and Art. 12 is formally mandatory for High Risk systems), Hive's Agent Transaction Graph exceeds the logging standard described in Article 12 as a matter of design.²

Agent Transaction Graph capabilities:

- Every agent action — contract execution, transaction initiation, dispute trigger, settlement confirmation — is appended as a cryptographically signed, immutable record to the Agent Transaction Graph
- Records include: agent DID, action type, timestamp (UTC), input parameters, output hash, and a chain link to the prior record
- No record can be deleted, modified, or reordered without invalidating the cryptographic chain — providing tamper evidence without a central audit authority
- Logs are queryable by enterprise auditors via authenticated API endpoint (see Section 10)

Mapping to Article 12 criteria:

Art. 12 Requirement	Hive Implementation
Automatic logging during operation	Agent Transaction Graph, appended at execution time
Level of traceability appropriate to intended purpose	Full action-level trace, not sample-based
Logging of input data where technically feasible	Input hashes logged; raw inputs available on request subject to data governance policy (Section 8)
Retention period	Configurable; default 90 days hot storage, indefinite cold archive

2. EU AI Act Art. 12 text: <https://eur-lex.europa.eu/eli/reg/2024/1689/oj> (Chapter III, Section 3)

4. Article 13 — Transparency Obligations

Art. 13 requires that high-risk AI systems be designed so that their operation is sufficiently transparent to enable deployers to interpret the system's output and use it appropriately.

Hive publishes the following transparency disclosures as standard:

1. System Card — available at thehiveryiq.com/docs/system-card — describes model architecture, training data provenance (not applicable — Hive is rule-based and contract-driven, not ML-trained), intended use, and known limitations
2. HAHS 1.0.0 contract schema — the machine-signed employment contract executed at agent spin-up is human-readable JSON-LD and is provided to any party that queries the agent's DID document
3. HiveTrust public risk scores — risk scores surfaced via GET https://hivetrust.onrender.com/v1/verify_agent_risk?agent_id=<DID> are accompanied by a `risk_factors` array that explains which inputs drove the score, satisfying the explainability requirement
4. SDK documentation — pip install `hive-civilization-sdk` ships with inline docstrings and a published API reference; no black-box behavior

Disclosure to natural persons: Where a HiveGate or HiveBank interaction is presented to a human end-user (e.g., a compliance officer reviewing a flagged transaction), Hive surfaces a plain-language explanation of the automated action taken and the agent responsible by DID.

5. Article 14 — Human Oversight

Art. 14 requires that high-risk AI systems be designed to allow natural persons to effectively oversee the system during its operation.

Hive's human oversight architecture includes the following controls:

Control	Description
HiveTrust override flag	Any agent's risk score can be manually overridden by an authorized administrator via authenticated API call; the override is itself logged to the Agent Transaction Graph
HAHS contract suspension	An authorized operator can suspend an agent's HAHS employment contract, halting all transaction authority, without requiring a code deployment
Transaction hold threshold	HiveBank transactions above a configurable USDC threshold are held for human approval before on-chain execution
Audit dashboard	Operators access a read-only audit view of the Agent Transaction Graph with full filter/export capability
Kill switch	Emergency agent deactivation endpoint available to credentialed operators; deactivation is logged and irreversible without re-onboarding

Hive's design philosophy is that agents operate with delegated authority, not autonomous authority — every consequential action is traceable to a human operator who provisioned the agent and defined its authority scope in the HAHS contract.

6. Article 22 — Automated Decision-Making

(GDPR Art. 22 Overlap)

EU AI Act obligations on automated decision-making interact with GDPR Article 22, which grants data subjects the right not to be subject to solely automated decisions that produce significant legal or similarly significant effects.³

HiveBank's automated transaction logic may, in downstream deployments, be applied to decisions that affect natural persons. Hive addresses this through:

explain_transaction endpoint:

GET /v1/explain_transaction?tx_id=<transaction_id>

This endpoint returns:

- The full decision chain that produced the transaction — agent DID, triggering condition, rule applied, timestamp
- The human-readable rationale in natural language
- The Agent Transaction Graph reference hash for the decision
- Flags if the decision was escalated to human review or executed automatically

This endpoint is designed specifically to satisfy:

- GDPR Art. 22(3) — the right to obtain human intervention, express one's point of view, and contest the decision
- EU AI Act Art. 13 — explainability of AI system output to deployers
- EU AI Act Recital 47 — meaningful information about automated logic

Important scoping note: GDPR Art. 22 obligations are borne by the data controller — typically the enterprise deploying HiveBank, not Hive Civilization as a processor. Hive's explain_transaction infrastructure enables data controllers to meet their own Art. 22 obligations without custom engineering.

3. GDPR Art. 22: <https://gdpr-info.eu/art-22-gdpr/>

7. Technical Robustness — HiveTrust Risk Verification

EU AI Act Art. 15 requires high-risk AI systems to achieve appropriate levels of accuracy, robustness, and cybersecurity. As a Limited Risk system, Hive implements equivalent controls as a matter of good engineering practice.

HiveTrust architecture:

- Risk scores are computed at query time from live agent state — they are not cached predictions; stale scores cannot be served
- Scores are deterministic given the same inputs — no probabilistic drift between equivalent queries
- The risk model is rule-based and auditable: weights and thresholds are documented in the HiveTrust technical spec (available to enterprise customers under NDA)
- The verification endpoint is rate-limited and authenticated to prevent enumeration attacks
- All traffic is TLS 1.3; DID resolution follows W3C DID Core specification

Adversarial robustness: Because Hive agents operate on-chain (Base L2), economic attacks must clear blockchain-level security properties. Agent identity cannot be spoofed without compromising the private key associated with the DID — a cryptographic guarantee, not a policy control.

Availability: HiveTrust is deployed on Render with automatic restart policies. Enterprise SLA terms are available on request.

8. Data Governance

8.1 Data Hive Stores

Data Category	Description	Storage Location
Agent DID documents	W3C DID Core compliant identity records	Distributed ledger + Hive registry
HAHS contract records	Machine-signed employment contracts	Hive secure database (encrypted at rest)
Agent Transaction Graph entries	Immutable action log records	Append-only log store
Risk score inputs	Agent behavioral signals used for HiveTrust scoring	Computed at query time; inputs not persistently stored beyond logging window
Natural person data (if any)	Only if enterprise deployer provisions human-linked agents	Data processor agreement required; Hive does not access

8.2 Data Minimization

Hive does not collect personal data about natural persons as part of its core agent network operations. Agent DIDs are not inherently linked to natural persons — they are identifiers for software agents. Any linkage between a DID and a natural person is established by the deployer, not by Hive, and governed by the deployer's privacy policy and GDPR obligations.

8.3 Retention

Data Type	Default Retention	Configurable?
Agent Transaction Graph	90 days hot / indefinite cold	Yes, per enterprise agreement
HAHS contracts	Duration of agent employment + 2 years	No (legal hold requirement)
Risk score query logs	30 days	Yes
DID documents	Active while agent is registered	Deleted on agent deregistration

8.4 Data Sovereignty

Enterprise customers may request geographic data residency constraints. Hive's default infrastructure is US-based. EU data residency is available as a contractual option for enterprise deployments targeting EU personal data.

8.5 Sub-processors

Hive uses the following sub-processors relevant to EU data flows:

- Render — compute and hosting (US-based)
- Base L2 (Coinbase) — blockchain settlement layer (globally distributed, on-chain data is public by design)

9. Self-Assessment Conclusion

Based on the analysis above, Hive Civilization's HiveBank and HiveGate services:

1. Do not fall within any prohibited AI practice under EU AI Act Art. 5
2. Are classified as Limited Risk under the current deployment scope, with obligations under Art. 13 (transparency) met through HiveTrust explain output, system card, and HAHS contract transparency
3. Exceed the logging and auditability standard of Art. 12 through the Agent Transaction Graph, which provides tamper-evident, cryptographically chained records of all agent actions
4. Support human oversight through configurable transaction holds, HAHS contract suspension, and the HiveTrust override mechanism
5. Enable downstream deployer compliance with GDPR Art. 22 and EU AI Act Art. 13 via the explain_transaction endpoint
6. Apply technical robustness controls — deterministic, rule-based risk scoring; on-chain identity security; TLS 1.3; rate limiting

Conformance statement: TheHiveryIQ declares, to the best of its knowledge and based on reasonable inquiry, that Hive Civilization's HiveBank and HiveGate services as described in this document conform to the applicable obligations of Regulation (EU) 2024/1689 (EU AI Act) for Limited Risk AI systems as of the date of this document.

This assessment will be reviewed and updated:

- Annually
- Upon material change to the Hive architecture
- Upon final publication of binding EU AI Act delegated acts or implementing regulations that affect this classification

Prepared by: Steve Rotzin, Founder, TheHiveryIQ

Contact: srotzin@me.com | thehiveryiq.com

Assessment date: April 2026

10. Live Endpoints for Auditors

The following endpoints are publicly accessible without authentication and may be used by enterprise auditors, procurement teams, and regulators to verify Hive's claims:

Endpoint	Purpose	Auth Required
GET /v1/verify_agent_risk?agent_id=<DID>	Real-time agent risk score + risk_factors explanation	No
GET /v1/explain_transaction?tx_id=<id>	Full decision chain for a HiveBank transaction	No (public transactions)
GET /v1/agent/<DID>/hahs	HAHS 1.0.0 contract record for a registered agent	No
GET /v1/agent/<DID>/transaction_graph	Agent Transaction Graph entries for a given DID	API key (enterprise)
PyPI: <code>pip install hive-civilization-sdk</code>	SDK for programmatic integration	No
thehiveryiq.com	System card, API reference, architecture docs	No

Enterprise auditors requiring access to authenticated endpoints or raw Transaction Graph exports should contact Steve Rotzin at srotzin@me.com to arrange a credentialed audit session.

This document is a draft self-assessment. It is not legal advice and does not constitute certification by a notified body. Enterprises should engage qualified EU AI Act counsel for deployment-specific compliance review.